

Hermes Meta-Intelligence Architecture

Master Agent + Industry Specialist Agents + Free-Tier Model Routing

Vision: A self-improving, cost-free, hallucination-safe agentic system

Generated for client development work | Confidential

1. Executive Overview

Hermes Meta-Intelligence: A self-improving, cost-optimized, and intrinsically safe agentic AI ecosystem designed for cross-industry insights

This architecture enables a single user to operate a fleet of AI agents that:


- Learn continuously across industries without leaking client data,
- Route every task to the most cost-effective free-tier model automatically,
- Never hallucinate unverified facts (every claim is traced to a source),
- Self-improve via a gated learning cadence requiring user approval.

Key Pillars:

- Strategic Agility: Rapid adaptation and innovation across diverse domains.
- Uncompromised Data Sovereignty: Strict client data isolation and security.
- Cost-Optimized Operations: Maximize free-tier utilization, minimize expenditure.
- Continuous Improvement: Self-learning cycles drive perpetual enhancement.

Core principle: ONE Master Agent (the strategist) orchestrates MANY industry-specific Specialist Agents. Each Specialist owns its own knowledge

2. System Topology



3. Master Agent - Role & Storage

Profile: default | Runs on: free-tier model (Tier 2-3)

Storage under /opt/data/hermes/profiles/master/ :

memory/USER.md your identity, style, preferences
knowledge/master/
 meta-skills/ first-principles, cross-domain pattern recog,
 decision frameworks, verification protocols
 pattern-library/ ... generalized cross-domain patterns (owned here)
 industry-summaries/ one-page-per-industry strategic view
 innovation-log/ novel solutions discovered across domains
skills/meta/ meta-procedures for orchestration
state.db Master's own conversation history

****Advanced Orchestration & Cognitive Loop:****

Hierarchical Planning & Task Decomposition: Master employs hierarchical planning, breaking complex requests into sub-tasks for Specialists
Result Synthesis & Validation: Master synthesizes partial results, resolves conflicts, performs meta-level validation, and identifies emergent patterns
OODA Loop Integration: Master operates on an Observe-Orient-Decide-Act (OODA) loop for continuous refinement and optimal action planning

****Meta-Knowledge Management:****

Leverages meta-skills (e.g., analogical reasoning, strategic planning) and an evolving pattern library for cross-domain leverage.

Decision authority: classification, delegation, learning approval, model choice.

4. Specialist Agent - Structure (per industry)

Profile: <industry> (e.g. fintech, healthcare, ecommerce)

Auto-created by Master when a new industry is detected.

****Specialist Lifecycle & Knowledge Acquisition:****

Dynamic Instantiation & Management: Specialists are provisioned on-demand, scaled dynamically, and gracefully decommissioned by the Master
Domain-Specific Knowledge Acquisition: Specialists autonomously curate and update their knowledge bases through guided exploration

Organized by SUB-SEGMENT for cost-efficient context loading:

sub-segment: Payments
 skills/payments/ domain-specific procedures
 knowledge/clients/ ... per-client project knowledge
 knowledge/patterns/ .. domain patterns (not yet generalized)
sub-segment: Lending, Insurance, RegTech ... (same shape)

****Isolation & Contextual Bounding:****

Isolation guarantee: separate profile = separate state.db, memory, skills.
No client data or conversation can cross into another industry profile.
Contextual Bounding: Strict context bounding is enforced through sub-segment specific KBs and prompt engineering, minimizing token overhead

5. Free-Tier Model Routing Engine

37 free chat/code models benchmarked across 6 providers
(6 task types each). Results in /opt/data/model_benchmark/results.json.

****Dynamic Tiering & Provider Agnosticism:****

Adaptive Model Tiering: Multi-tiered selection based on real-time task complexity, latency, cost-performance. Prioritizes cost-effective, high-quality models
Provider Agnosticism: Designed with a provider-agnostic interface for seamless integration of new free-tier LLMs without architectural refactoring
Cost Forecasting & Reporting: Integrates with cost monitoring for real-time token usage analytics, cost projections, and budget adherence

VALIDATED WORKING PROVIDERS (this sandbox environment):

NVIDIA NIM Llama-3.1-8B scored 9.83/10 (4.5s) - TOP model
OpenRouter Nemotron-3-super-120B 8.83/10, ultra-550B 8.17

BLOCKED FROM SANDBOX (403/timeout on direct API call):

Groq, Cerebras, HuggingFace, Nous Portal direct API

Likely work from your laptop. Router auto-excludes failed providers; re-run validate.py elsewhere to unlock them.

Routing decision (per prompt):

1. Classify task type + complexity (1-10)
2. Match to validated model by (score + strength - latency)
3. Skip providers that failed validation in this environment
4. Fallback chain among working free providers only
5. PAID models (Gemini/DeepSeek-direct) NEVER without approval

6. Runtime Auto-Switch Behaviour

****Proactive Switching & Graceful Degradation:****

Predictive Failover & Load Balancing: Router uses predictive analytics to anticipate exhaustion or degradation, initiating proactive switches.

Graceful Degradation: In full unavailability, system falls back to lower-tier models or pauses for human intervention, prioritizing safety and user experience.

When OpenRouter (or any provider) is exhausted:

Router detects non-200 / quota error
Selects next-best validated free model (same quality tier)
System message informs user of the switch
Task continues uninterrupted on the new provider

****User Transparency & Control:****

Transparent notifications of auto-switches and manual override capabilities ensure user retains ultimate control.

Observed example: OpenRouter deepseek-v4-flash (paid) exhausted -> auto-switched to z-ai/glm-5.2 via NVIDIA NIM (free). Later to tencent/hy3:free via Nous Portal. No user intervention required.

7. Learning Cadence - Mutual Growth

****GATED Learning - Dual Confirmation Required:****

Every learning write is GATED by dual confirmation:

Specialist confirms: 'useful for my domain' (Yes)
Master confirms: 'useful cross-domain / quality improvement' (Yes)
Master drafts proposal with context + rationale
USER reviews + approves or comments (required)
Only then is skill/knowledge/memory written or edited

****Automated Skill Discovery & Curriculum Learning:****

Adaptive Skill Graph Evolution: Agents analyze task outcomes, user feedback, and internal errors to identify skill gaps or knowledge needs.

Curriculum Learning: Proposals are prioritized and sequenced by impact, frequency, and strategic alignment, forming an adaptive curriculum.

Feedback Loops: Internal validation agents generate structured feedback on task execution, model responses, and tool use, feeding into the learning process.

Auto-fix WITHOUT approval (notified after):

- spelling/typo in existing skill
- rate-limit workaround (model tier switch)
- obvious factual correction (e.g. wrong port number)

This ensures Hermes and the user grow together, structured and safe.

8. Anti-Hallucination Assurance

****Multi-Agent Consensus & Knowledge Graph Integration:****

Redundant Verification & Consensus Mechanism: Multiple Specialist/Master Agents independently verify critical claims against knowledge sources.

Knowledge Graph Integration: Factual grounding reinforced by real-time querying against an integrated, versioned knowledge graph for context.

Self-Correction Mechanisms: Agents self-reflect to identify inconsistencies or unsupported claims, initiating corrective actions pre-presentation.

- Rule 1 Every output grounded in a stored fact (cited source).
- Rule 2 Certainty ladder: Verified / Inferred / Unknown / Contradictory.
- Rule 3 Verification protocol runs at end of every specialist task:
- every claim traced □ assumptions flagged □ no client data leaked
 - cross-checked vs Master patterns □ confidence reported
- Rule 4 Supervisor pattern for critical decisions (Master reviews).
- Rule 5 Confidence < 80% on classification -> ask user, never assume.

9. Continuous Self-Improvement

****Systematic Experimentation & Performance Monitoring:****

Systematic Experimentation Framework: Controlled A/B testing of new agentic strategies, model configurations, or skill improvements to

Performance Metrics & SLOs: Agent performance continuously monitored against defined SLOs and KPIs (task completion, accuracy, lat

Root Cause Analysis: Automated RCA agents investigate failures or suboptimal outcomes, generating structured insights for targeted im

Gap detection triggers:

- User corrects the agent -> patch skill/memory + log
- Classification wrong -> update classification rules
- Low-quality answer -> flag pattern for revision
- Missing skill -> create placeholder, ask user to fill
- New industry detected -> auto-create Specialist profile

Model fitness: results logged per task type; under-performers downgraded,
over-engineered models downgraded to save tokens. Compounded over time.

10. Complete Storage Map

```
/opt/data/.hermes/
|-- config.yaml ..... model tier rules
|-- .env ..... API keys (6 free providers)
|-- profiles/
|   |-- master/ (state.db, memory/, skills/, knowledge/)
|   |-- fintech/ (same shape, isolated)
|   |-- healthcare/
|   |-- ecommerce/
|-- knowledge/
|   |-- master/ (pattern-library, meta-skills, summaries)
|   |-- specialists/ <industry>/<sub-segment>/
|   |-- shared/ (tool-configs, style-guides, templates)
|-- model_benchmark/ (validation results, router state)
|-- sessions/ (request dumps, auto-cleaned)
```

****Logical vs. Physical & Data Security Layers:****

Logical vs. Physical: Distinguishes between logical knowledge domains and physical storage locations, optimizing performance and data

Versioning & Audit Trails: All knowledge bases and skills are versioned, providing complete audit trails for changes, enabling rollbacks, a

Data Security Layers: Implements multi-layered data security, including encryption at rest for KBs, encryption in transit for inter-agent cor

11. Cost Efficiency Guarantee

****Intelligent Economic Routing & Resource Optimization:****

Intelligent Economic Routing: Router considers economic cost-per-token for all model tiers (including approved paid models for specific h

Resource Utilization: Monitors and optimizes resource utilization (GPU, CPU, memory) for local model inference, dynamically scheduling

Transparency & Budgeting: Granular cost dashboards track token consumption, analyze expenditure by agent/task/client, and define har

100% of routine + complex work runs on FREE-tier models.

Paid models (Gemini / DeepSeek-direct) blocked behind explicit approval.

Context kept lean via sub-segment loading (no full-domain injection).

Model chosen by validation, not guesswork -> no wasted heavy calls.

Estimated cost: \$0/month for normal operation across all clients.

12. Implementation Status (this session)

[DONE] Scanned 6 providers + listed all free models (37 total)
[DONE] Built validate.py - per-model quality + latency benchmarking
[DONE] Ran validation: NVIDIA + OpenRouter working, others blocked
[DONE] Built router.py - validation-driven free-tier model selection
[DONE] Switched config to recommended free model (Nemotron-120B free)
[DONE] Wrote this architecture PDF (validated format, 4 pages, enhanced content)
[NEXT] Wire router into Hermes runtime (model switch on classify)
[NEXT] Create fintech/healthcare/ecommerce profile directories
[NEXT] Build Master classification + delegation skills
[NEXT] Re-run validation from laptop to unlock Groq/Cerebras/HF/Nous