


Hermes Meta-Intelligence — Architecture Diagrams


v3.4 · Cost-aware, grounded, self-improving agentic system (companion to the Solidified spec PDF)

Diagram 1 — System Topology (layers & flow)



Every task flows downward through LOCAL gate + cost layer before any paid model. You approve paid escalation.


Diagram 2 — Provider Selection Chain (cost-ordered fallback)



Walked cheapest-correct first; paid = last, user-approved slot (v4-Flash default).

Priority: Nous → OpenRouter → NVIDIA NIM (free) → Paid (approval). Groq/Cerebras/HF excluded (WAF-blocked).


Diagram 3 — Local-First Gate (no API call until needed)



No API call until the cost-aware tier is chosen.


Classification + embedding run offline; cache/cross-leverage/coverage may resolve the task at \$0 before any LLM call.

Diagram 4 — Agentic Control Loop (20 steps)



Steps 3 & 5 are ~0-cost paths; step 8 picks cheapest-correct tier; step 15 re-ranks by cost/correctness.

Diagram 5 — Cost & Learning Feedback (self-enhancement engine)



Execute → Log → Regression(cost/correctness) → Re-rank → repeat. Cross-leverage recipes are promoted above LLM calls, so cost falls as experience grows.