

OAKAI Mission & Roadmap

Enterprise AI Business Solution Provider — Executive Brief

PREPARED FOR	Weng Guan / Founder, OAKAI SDN BHD
REFERENCE	OAKAI-MIS-001
DATE	2026-08-13
CLASSIFICATION	OAKAI Confidential

Table of Contents

01	Executive Summary	3
02	MVP & Market	4
03	Business Model	5
04	COO Roadmap	7
05	Technology Foundation	9
06	Governance & Controls	11
07	Financial Snapshot	13

01

Executive Summary

Mission, positioning, and current state

OAKAI SDN BHD is building an AI-powered business solution provider that targets operational excellence gaps across Manufacturing, Systems Integration, and Administration/Back-Office sectors. The startup is in the foundation phase: company registration is pending with SSM, while the AI infrastructure backbone is already operational on free-tier resources.

Key Positioning

AI solutions backed by a mental-model approach — not just tool integration

Free-tier-first economics (Nous Portal, OpenRouter, local LLM) eliminates upfront infrastructure cost

Vertical overlap: expertise in manufacturing/power-house, systems integration, and back-office creates a unique intersection advantage for cross-industry AI solutions

POC-first methodology: validate with real customer pain before scaling to enterprise contracts

The company operates on a lean bootstrap model with a hard cap on paid API usage (founder approval required). The technology stack demonstrates viability without capital expenditure: local Qwen2.5-1.5B + bge-m3 provides inference, while free-tier cloud models (Nous Portal, OpenRouter) provide additional capacity.

02

MVP & Market

Minimum viable product + target positioning

The MVP stack is built on free-tier + local infrastructure, proving the business model works without paid compute:

Component	What It Is	Current Status	Owner
Local AI Stack	Qwen2.5-1.5B + bge-m3 inference on local machine	■ Running	Ops
COO Brief Automation	Weekly strategic guidance generated via cron	■ Week 1 delivered	AI Co-pilot
PENSOLAR POC	Solar PV project management with AI anomaly detection	■ Scaffolded	AI Co-pilot + Founder
Mentor System	Daily AI education in business language	■ 3 concepts delivered	AI Mentor
Marketing System	LinkedIn presence + daily content cadence	■ In progress	Founder + cron
GitHub Presence	Open-source knowledge + skills repo	■ Active (5 commits)	AI Co-pilot

Market positioning leverages three competitive advantages:

Advantage	Description	Differentiator
Mental-Model Approach	Solutions grounded in systematic problem-framing, not just tool assembly	Most AI consultants skip the 'why'; we teach the 'how to think'
Free-Tier Economics	Full operational stack on free resources — zero capex infrastructure	Competitors spend on cloud; we prove ROI before asking clients to
Vertical Overlap	Manufacturing + SI + Admin expertise creates unique cross-vertical solutions	Single-vertical consultants miss 60% of integration opportunities

03

Business Model

VTDF framework: Value, Tech, Distribution, Finance

Adapted from the VTDF (Value-Technology-Distribution-Finance) framework for AI consulting services. Unlike pure SaaS plays (e.g. C3.ai), OAKAI is positioned as a solution provider: the product is implemented AI systems, not licensed software.

03A

Value Proposition

What clients buy

Value Driver	Client Benefit	Evidence
Operational Excellence	Reduce waste, variation, and delay in core workflows	7 Pillars framework validated in PENSOLAR case
AI-Augmented Decision Making	Real-time exception detection + recommendation engine	Qwen2.5-1.5B serving anomaly-flagging demo for solar PM
Risk-Gated Autonomy	AI systems that can act, but stop at the right risk threshold	Shadow→Canary→Enforce rollout with recall-gated auto-rollback
Evals-Based Trust	Every AI capability measured against frozen golden sets	golden_v1.csv + score_eval.py: recall gate on guardrails

03B

Technology Stack

The AI backbone stack

Layer	Component	Provider	Cost Tier
LLM (brains)	poolside/laguna-s-2.1:free	Nous Portal	Free (~50RPM)
LLM (fallback)	tencent/hy3:free	Nous Portal	Free (cron-pinned)
Local LLM	Qwen2.5-1.5B + bge-m3	Local Ollama	Free (zero egress)
Routing	Free-tier model selection policy	Custom (hermes-model-tiering)	Free
Eval Harness	Precision/recall + golden sets	Custom (score_eval.py)	Free
Guardrails	Approval gate + input filtering	Hermes model-selection-policy	Free

Layer	Component	Provider	Cost Tier
Orchestration	Hermes Agent + cron jobs	Local s6 + gateway	Free
Storage	SQLite state.db + git version control	Local + GitHub	Free

03C

Revenue Model

Path to paid services

Stream	Model	When
AI Consulting (POC)	Fee-for-deliverable: POC builds, eval design, guardrail setup	Weeks 1-8
AI Consulting (Enterprise)	Retainer + success-based: deployed agents, training, audit	After 2 client POCs
Knowledge Products	Templates, frameworks, golden-set scaffolding sold as packages	Q1 2027
Training Programs	Workshop: mental-model AI for operations teams	Q2 2027

Unit economics target: LTV:CAC > 3:1, gross margin > 70%, payback period < 90 days. First 3 client POCs will be at reduced rate to build case studies.

04

COO Roadmap

Week-by-week execution plan

The COO guidance (produced by the *strategic-coo-guidance* cron every Sunday 08:00 MYT) frames execution in 3-stream parallel sprints. Each week has a legal/marketing/product stream, with success metrics tied to measurable outcomes.

Week	Legal Stream	Marketing Stream	Product Stream	Key Metric
W1	SSM incorporation + domain	LinkedIn profile + Day-1 post	3 client surveys + local LLM scaffold	SSM receipt + 30 LinkedIn followers
W2	Bank account opening	Landing page GitHub Pages	PENSOLAR demo scenario 1	Bank ready + landing page live
W3	Service agreement template	Content calendar + 7 posts	PENSOLAR demo scenario 2	3 demo scenarios drafted
W4	NDA template + COO Brief package	Group engagement 15+ posts	Local LLM integration validated	First client demo tested
W5-W8	Client contract templates	Case study content pipeline	Client POC #1 (real engagement)	First paid POC signed

04A

Risk Mitigation

Proactive threat coverage

Risk	Impact	Probability	Mitigation
Egress blocked (HF DNS, Groq WAF)	H	H	Local-first stack: Qwen2.5-1.5B + bge-m3
Free-tier throttled (Nous ~50RPM, OR 50/day)	M	H	Batch off-line; local LLM is primary
SSM name rejection	H	M	3 names filed simultaneously
No client responds to survey	M	M	Lead with personal network + RM20 incentive each
LinkedIn shadow-ban	L	M	Human-first posting, spaced cadence, no hashtag spam
Bank account delay	H	M	Prep packet day 4; founder personal bridge account tagged

04B

Success Metrics

Week 1 targets (gate: pass/fail)

Metric	Baseline	Target	Signal
Company status	Unnamed	SSM receipt issued	Green = A1+A2 started
Domain	Unregistered	Paid + DNS configured	Green = A3 done
Bank account	None	Prep packet complete	Green = A4 done; open = W2
LinkedIn followers	0	30+	Green ≥ 30
Client surveys	0	3 (mfg/retail/F&B;)	Green = C1 done
Local LLM scaffold	None	Running inference	Green = C5 done
W1 cash spent	—	≤ RM215 core	Budget held

Pass threshold: A1+A2+A3 landed, C5 scaffold running, 1 survey done, Day-1 post live, 30+ followers.

05

Technology Foundation

The AI backbone in practice

The OAKAI AI stack consists of five layers, each chosen for cost-free operation while maintaining production viability:

Layer	Component	Function
LLM Layer	Qwen2.5-1.5B (local) + laguna-s-2.1 (Nous free)	Reasoning, synthesis, copy drafting
RAG Layer	bge-m3 embeddings + SQLite FTS5	Knowledge retrieval, document search
Agent Layer	Hermes Agent + skills system	Autonomous task execution, tool use
Guardrail Layer	model-selection-policy skill	Free-first routing, approval gates, cost control
Eval Layer	golden_v1.csv + score_eval.py	Precision/recall, recall-gated rollout

Key architectural decisions:

Decision	Choice	Rationale
Free-tier priority	Nous Portal → OpenRouter → NVIDIA NIM → Paid (approval)	Minimizes cash burn; paid requires founder sign-off
Local fallback	Qwen2.5-1.5B + bge-m3 via Ollama	Egress blocked (HF DNS, Groq WAF); local-first mandatory
Model pinning	cron jobs pinned to tencent/hy3:free	Prevents model-drift from default changes breaking crons
State protection	state.db gitignored + pre-commit hook blocks commits	Prevents Aug 12 race-condition corruption
Gateway	s6-managed gateway service	Auto-restart on crash; notifications via messaging platforms

05A

Autonomous Systems

Six parallel cron streams

Cron Job	Schedule	Deliverable
mentor-ai-daily	15:00 MYT x3 (07/15/22)	AI concept + test in knowledge/mentor/

Cron Job	Schedule	Deliverable
learn-pensolar	15:00 MYT	Gap-driven research update to PENSOLAR SUMMARY
strategic-coo-guidance	Sun 08:00 MYT	Weekly COO brief in knowledge/strategy/
marketing-advisor-daily	06:00 MYT Mon-Sat	Daily marketing content for LinkedIn
workspace-cleanup-daily	02:00 MYT	14-day retention pruning
startup-catchup-enforcement	06:00 MYT daily	Auto-enforce startup checks + catchup

06

Governance & Controls

How safety and quality are enforced

Three layers of defense ensure quality output and safe operation. These are not aspirational — they are implemented and verified in the current stack:

06A

Approval Gate

Before paid API calls

The *model-selection-policy* skill enforces free-first routing. Paid models (Gemini, DeepSeek direct) require explicit founder approval. The routing chain: Nous Portal → OpenRouter → NVIDIA NIM → Paid (approval-gated).

06B

Eval Gate

Before trusting AI output

Every AI tool is evaluated against frozen, stratified golden sets (golden_v1.csv, 20 rows, 25% block cases). Precision/recall/accuracy reported. Recall is the gate for guardrails — accuracy is the last number reported, not the first.

06C

Autonomy Gate

Safe AI action in production

Shadow → Canary → Enforce rollout: shadow (advisory, human decides) → canary (low-risk auto-approval) → enforce (full autonomy). Auto-rollback to dial 0 if recall drops below threshold on any block case.

06D

Weekly Compliance Checklist

Must-pass before closing the week

- state.db healthy (sessions > 0, integrity check PASS)
- Gateway service running (s6 supervision, auto-restart enabled)
- All cron jobs executed (no missed runs > 1 cycle)
- No paid API calls without founder approval (check model_config_snippet.yaml)
- Golden set audit trail updated (if any new cases labeled)
- GitHub push verified (local HEAD == remote, no secrets committed)

- Free-tier quotas not exhausted (Nous < 50RPM, OR < 50/day)
- Egress blockers confirmed (HF DNS still blocked, Groq/WAF still active)

07

Financial Snapshot

Bootstrap economics + W1 budget

Item	Cost (RM)	Tier	Status
SSM name reservation	30	Core	Pending
SSM incorporation (private limited)	110	Core	Pending
Domain oakai.com.my (1yr)	15	Core	Pending
Survey incentive (3 clients)	60	Core	Pending
Subtotal core	215		
Contingency buffer	45	Optional	Pending
LinkedIn Premium (W2)	—	Deferred	Deferred
Paid APIs (W1+)	—	Approval-gated	Not used
W1 cash max	260	Core + buffer	Cap

Key Constraint

Egress reality: HF DNS-blocked, Groq/Cerebras behind Cloudflare WAF. All dev/POC runs on local stack. Cloud AI only via approved free tiers.

Free-tier AI (Nous Portal ~50RPM, OpenRouter 50/day) + local Qwen2.5-1.5B cover all W1 synthesis, copy, and contract drafting. No paid LLM call is made without founder approval.